

Optimus is here!

Self taught game playing AI from pixels

BY

- ***Raja Ayyanar***
- ***Jhansi***
- ***Divya***

Introduction

- Artificial intelligence agent playing games.
- How to make a machine Intelligent?
 - Supervised, Unsupervised and Reinforcement
- How smart AI is, as of today in gaming field?
 - Smart enough to beat any human in playing 2D games
 - Not to forget Google's AlphaGO Zero which demolished stockfish and komado in Chess.
 - Learns just like Human from Experience

Can we use Supervised Learning?

- Need Large datasets of games played.
- Search space is very huge.
- Will work for only onegame at a time.
- Need different architecture for different games.
- Impossible to make because of overfitting in DNN.

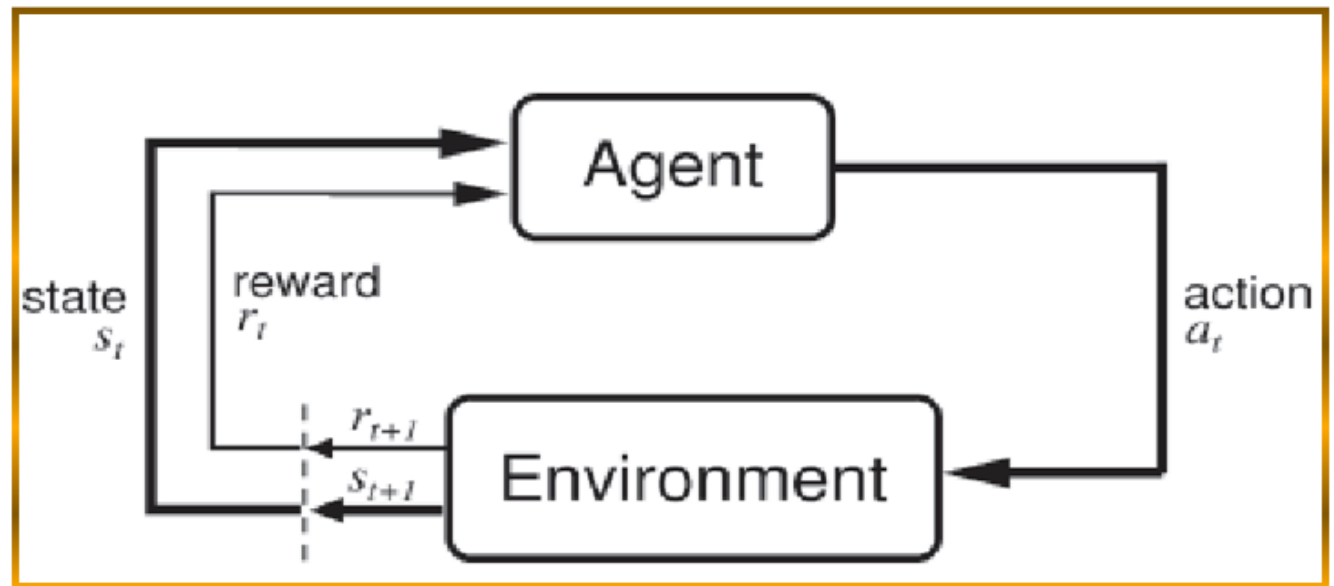
Hardness of problem

- State action π policy can not be optimized by supervised learning.

Deep Reinforcement Learning

- Definition:
 - branch of Machine learning which learns from trial and error. Based on experience and rewards.
- Markov Decision Process
 - Mathematical formulation of reinforcement learning scenario

Sets of state,
Sets of Actions
Reward function
Policy π
Value function, V



Design: Deep Q Learning

- Q function is a function which predicts action for a particular state

Algorithm 1 Deep Q-learning with Experience Replay

Initialize replay memory \mathcal{D} to capacity N

Initialize action-value function Q with random weights

for episode = 1, M **do**

 Initialise sequence $s_1 = \{x_1\}$ and preprocessed sequenced $\phi_1 = \phi(s_1)$

for $t = 1, T$ **do**

 With probability ϵ select a random action a_t

 otherwise select $a_t = \max_a Q^*(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in \mathcal{D}

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from \mathcal{D}

 Set $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ according to equation 3

end for

end for

| S.No | Content | Values | description |
|------|------------------------|--|---|
| 1 | Architecture | 6 layer DNN | |
| 2 | Input | State values from game simulator | Game Pixels or pre-processed info from pixel data |
| 3 | Output | Probability percentage of each actions | Actions are the possible variables that can be controlled for a particular game |
| 4 | Hidden Layer 1 | 160 neurons, Relu Activation | Hyper parameter |
| 5 | Hidden Layer 2 | 250 neurons, LRelu activation | Hyper parameter |
| 6 | Hidden Layer 3 | 500 neurons, Relu Activation | Hyper parameter |
| 7 | Hidden Layer 4 | 250 neurons, LRelu Activation | Hyper parameter |
| 8 | Hidden Layer 5 | 100 neurons, Relu Activation | Hyper parameter |
| 9 | Output Layer | No of actions as no. of neurons, Softmax layer | Number of possible actions is same as neuron output |
| 10 | Optimizer | ADAM | Generalized Adagrad |
| 11 | Learning Rate | 0.001 | Hyper parameter |
| 12 | Loss function | Cross Entropy loss | Log loss will be differentiable |
| 13 | Optimization technique | Back prop | Deterministic approach |
| 14 | Regularization | Dropout | To avoid over fitting |

Simulation: Game Environment

- Using OpenAI's Gym Library
- Pixels as a input
- Convolution Neural Network can be used if we use without preprocessing
- Gym creates an game environment which takes an action for particular state and gives reward, new state and info.

<https://github.com/openai/gym/wiki/Table-of-environments>

Total of more than **80 different Game** environment is available. Our code will run for all games with small tweaks.

Experimentation

- Results Obtained: we chose easiest game for demo(To save training time)

Training: For cartpole game

| No of trial games | Mean score | % of prediction correctly | % of action 1 is chosen | % of action 0 is chosen |
|-------------------|------------|---------------------------|-------------------------|-------------------------|
| 10 | 9 | 4.3% | 70 | 30 |
| 100 | 14 | 4.1% | 36 | 64 |
| 500 | 17 | 7.00% | 69 | 31 |
| 1000 | 30 | 10.56% | 60 | 40 |
| 5000 | 190 | 69.47% | 55 | 45 |

Performance

- Testing Result
 - After Learning for 5000 trials.

| s.No | No.of test games | Mean score | Max score | % of action 1 chosen | % of action 0 chosen |
|----------------|------------------|------------|-----------|----------------------|----------------------|
| Validation n 1 | 10 | 254 | 390 | 48.46 | 51.53 |
| Validation n 2 | 10 | 188 | 368 | 52.13 | 47.87 |
| Validation n 3 | 10 | 289 | 480 | 50.88 | 49.11 |

Demonstraion

- Tools used:

| S.No | Content | Value |
|------|----------------------|--|
| 1 | Programming Language | PYTHON |
| 2 | Game Environment | GYM library |
| 3 | Deep Neural Network | Tf-layer, Keras with Tensorflow Backend |
| 4 | Operating System | UBUNTU (because gym works better with LINUX) |
| 5 | Environment | Docker, Anaconda |
| 6 | IDE | Spyder, Jupyter |

Application

- Generalised AI for all 2D games
- Self learning Helicopter flying
- Self taught Car drifting
- Self taught crawling robots
- Finance Market- to learn optimal trading strategy
- Manufacturing sector- Robot to pick tools and arrange
- Robot learning

Potential for commercialization

- Deep Reinforcement Learning:

Combination of Deep Learning + Reinforcement learning has led to:

1. AlphaGo beating world champion at GO
2. Self-driving Cars (Tesla)
3. Video-game playing machines that plays at superhuman level

RL has been around since the 70s But proof of its greatness has only been recently demonstrated

Budget

- To build game playing AI – Zero
- Since most software and libraries are open source.
- But a GPU availability will be useful. Since most of the games took around 10 to 36 hours for training in a CPUs.

Findings and Conclusion

- 1.The Potential of DQL and DDQL for Reinforcement learning to build an more intelligent machines are yet to come.
- 2.RL is on bleeding edge of what we can do with AI.
- 3.As Professor Andrew NG said, "Artificial Intelligence is new electricity". We are near to see AI everywhere soon.



Credits

- Paper from Google DeepMind's Deep-Q-learning algorithm
- Andej Karpathy's blog
- CS231n classroom lecture from stanform university
- Keon.io for demo/Explanation on RL
- Sentdex youtube channel for lecture videos on coding RL and python libraries.